

# DEPLOYING A HYBRID INFRASTRUCTURE FOR RESEARCHERS IN AWS

Ihona Maria Correa de Cabo

H O L O 2

High-performance and distributed  
computing for big data

2023-2024

# INDEX

<b>1. INTRODUCTION</b>	<b>1</b>
<b>2. INFRASTRUCTURE SET UP</b>	<b>2</b>
2.1. VPC and subnets	2
2.2. Route tables	4
2.3. Security groups	6
2.4. EC2 instances	8
2.5. S3 buckets	11
<b>3. OPEN VPN CONFIGURATION</b>	<b>11</b>
<b>4. JUPYTER NOTEBOOK SETUP</b>	<b>14</b>
<b>5. VOILA SERVER CONFIGURATION</b>	<b>18</b>
<b>6. CHALLENGES AND TROUBLESHOOTING</b>	<b>19</b>
<b>7. CONCLUSIONS</b>	<b>21</b>

# 1. INTRODUCTION

The aim of this assignment is to create a hybrid cloud infrastructure within Amazon Web Services (AWS) to support a research team that uses Jupyter Notebooks for data analysis. This setup will involve a mix of public and private resources to allow secure analysis of data while providing a way to share findings with the public.

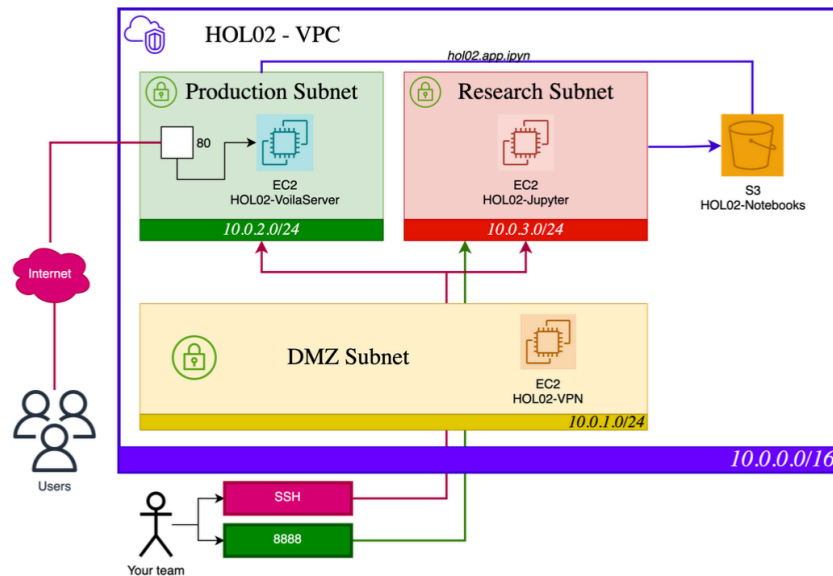


Figure 1.- Architecture of the proposed hybrid infrastructure.

This architecture contains a Virtual Private Cloud named HOL02 - VPC, which is an isolated cloud environment within AWS that holds all the resources. Inside this VPC, there are 3 subnets:

- Production Subnet: Contains the HOL02-VoilaServer EC2 instance, which is exposed to the internet via port 80 (HTTP). Ideally, by using the Voila Server, which turns Jupyter Notebooks into web applications, the team can share their results publicly without exposing the data or analysis environment.
- Research Subnet: Contains the HOL02-Jupyter EC2 instance, which hosts the Jupyter Notebook server. It is private and can only be accessible through the VPN.
- DMZ (De-Militarized Zone) Subnet: Contains the HOL02-VPN EC2 instance. It is a semi-secure area that provides a buffer between the untrusted internet and the protected internal network. This subnet hosts the VPN server that the team can use to securely connect to the research subnet.

The VPC is connected to the internet via an internet gateway named HOL02-IGW, allowing the HOL02-VoilaServer to share content to the public internet.

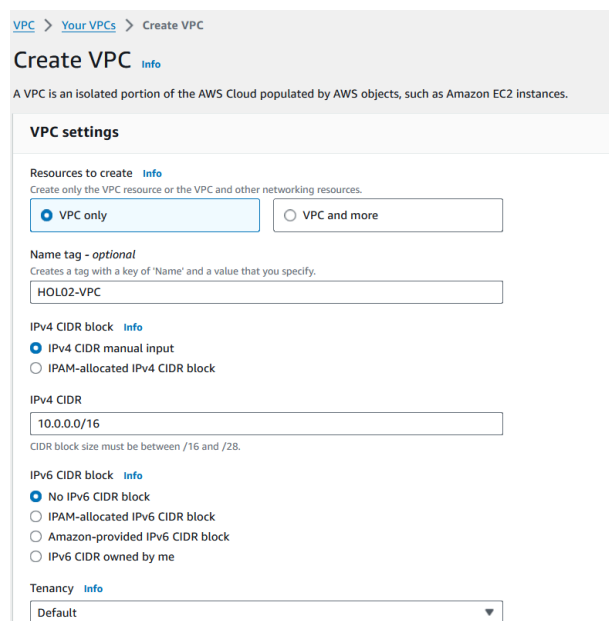
Each server (VPN, VoilaServer, Jupyter) is an EC2 instance, which is a virtual server in AWS's cloud.

The architecture is designed to provide a secure environment for the research team to analyze data using Jupyter Notebooks while at the same time ensuring that the results can be shared with the public through a web server, without compromising the research environment's security.

## 2. INFRAESTRUCTURE SET UP

### 2.1. VPC and subnets

First of all, the VPC is created. In order to achieve that, we must log into the AWS management console and navigate to the VPC dashboard. Then, we create the VPC named "HOL02-VPC". The "IPv4 CIDR block" field is set to "10.0.0.0/16".



VPC > Your VPCs > Create VPC

### Create VPC [Info](#)

A VPC is an isolated portion of the AWS Cloud populated by AWS objects, such as Amazon EC2 instances.

#### VPC settings

**Resources to create** [Info](#)  
Create only the VPC resource or the VPC and other networking resources.

☒ VPC only ☐ VPC and more

**Name tag - optional** [Info](#)  
Creates a tag with a key of 'Name' and a value that you specify.

HOL02-VPC

**IPv4 CIDR block** [Info](#)  
☒ IPv4 CIDR manual input  
☐ IPAM-allocated IPv4 CIDR block

IPv4 CIDR  
10.0.0.0/16  
CIDR block size must be between /16 and /28.

**IPv6 CIDR block** [Info](#)  
☒ No IPv6 CIDR block  
☐ IPAM-allocated IPv6 CIDR block  
☐ Amazon-provided IPv6 CIDR block  
☐ IPv6 CIDR owned by me

**Tenancy** [Info](#)  
Default

Figure 2.- VPC creation.

Next, from the VPC Dashboard, we select "Subnets" in the sidebar.

A subnet is a segmented portion of a VPC that allows to partition this VPC into smaller networks to better organize and secure the resources.

According to the assignment, we must create 3 subnets, one for the production, another for the research group, and finally the DMZ subnet.

When creating the subnets, we select the newly created HOL02-VPC from the drop menu and then create each of the subnets in the same way as *Figure 3*, each subnet with a different IPv4 subnet CIDR block.

**Create subnet** [Info](#)

**VPC**

VPC ID  
Create subnets in this VPC.  
vpc-0399cd2bbc16bd746 (HOL02-VPC)

Associated VPC CIDRs  
IPv4 CIDRs  
10.0.0.0/16

**Subnet settings**  
Specify the CIDR blocks and Availability Zone for the subnet.

**Subnet 1 of 1**

Subnet name  
Create a tag with a key of 'Name' and a value that you specify.  
HOL02-DMZ  
The name can be up to 256 characters long.

Availability Zone [Info](#)  
Choose the zone in which your subnet will reside, or let Amazon choose one for you.  
No preference

IPv4 VPC CIDR block [Info](#)  
Choose the VPC's IPv4 CIDR block for the subnet. The subnet's IPv4 CIDR must lie within this block.  
10.0.0.0/16

IPv4 subnet CIDR block  
10.0.1.0/24 256 IPs

Figure 3.-DMZ subnet creation.

**Subnets (9)** [Info](#)

Find resources by attribute or tag

<input type="checkbox"/>	Name	Subnet ID	State	VPC	IPv4 CIDR
<input type="checkbox"/>	HOL02-DMZ	<a href="#">subnet-08dbf0a9677464b08</a>	Available	<a href="#">vpc-0399cd2bbc16bd746</a>	10.0.1.0/24
<input type="checkbox"/>	-	<a href="#">subnet-0065246d17372ef01</a>	Available	<a href="#">vpc-0ce649878ba820244</a>	172.31.0.0/20
<input type="checkbox"/>	-	<a href="#">subnet-0bc0075288d7efcc5</a>	Available	<a href="#">vpc-0ce649878ba820244</a>	172.31.64.0/20
<input type="checkbox"/>	-	<a href="#">subnet-085274192e19b9165</a>	Available	<a href="#">vpc-0ce649878ba820244</a>	172.31.32.0/20
<input type="checkbox"/>	-	<a href="#">subnet-0598bcf1a085db18e</a>	Available	<a href="#">vpc-0ce649878ba820244</a>	172.31.80.0/20
<input type="checkbox"/>	-	<a href="#">subnet-07a4580da069fa397</a>	Available	<a href="#">vpc-0ce649878ba820244</a>	172.31.16.0/20
<input type="checkbox"/>	-	<a href="#">subnet-0f8bca81519fb3d54</a>	Available	<a href="#">vpc-0ce649878ba820244</a>	172.31.48.0/20
<input type="checkbox"/>	HOL02-Production	<a href="#">subnet-0e0f5e298977dbc58</a>	Available	<a href="#">vpc-0399cd2bbc16bd746</a>	10.0.2.0/24
<input type="checkbox"/>	HOL02-Research	<a href="#">subnet-05430e9d59aeecca</a>	Available	<a href="#">vpc-0399cd2bbc16bd746</a>	10.0.3.0/24

Figure 4.- All the subnets created.

Each CIDR block represents a range of IP addresses that can be assigned to resources within that subnet. By assigning different CIDR blocks to different subnets,

we ensure that the IP addresses used in one subnet do not overlap with the ones used in another subnet within the same VPC.

## 2.2. Route tables

In AWS, a route table is a set of rules, called routes, that are used to determine where network traffic from the VPC should be directed. Each subnet in a VPC must be associated with a route table, and the route table controls the routing of traffic within that subnet.

To proceed, we go back to the VPC Dashboard and select "Route Tables" from the sidebar. We create one route table per each subnet:

VPC > Route tables > Create route table

### Create route table [Info](#)

A route table specifies how packets are forwarded between the subnets within your VPC, the internet, and your VPN connection.

**Route table settings**

**Name - optional**  
Create a tag with a key of 'Name' and a value that you specify.

HOL02-DMZ-RT

**VPC**  
The VPC to use for this route table.

vpc-0399cd2bbc16bd746 (HOL02-VPC)

**Tags**  
A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

Key Value - optional

Q Name X Q HOL02-DMZ-RT X Remove

Add new tag

You can add 49 more tags.

Cancel Create route table

Figure 5

Route tables (5) [Info](#)

Find resources by attribute or tag

<input type="checkbox"/>	Name	Route table ID	Explicit subnet associ...	Edge associations	Main	VPC	Owner ID
<input type="checkbox"/>	-	<a href="#">rtb-0289d69f294ac3cb8</a>	-	-	Yes	<a href="#">vpc-0ce649878ba820244</a>	654654181466
<input type="checkbox"/>	-	<a href="#">rtb-033cc6751c1a9394d</a>	-	-	Yes	<a href="#">vpc-0399cd2bbc16bd746   HOL...</a>	654654181466
<input type="checkbox"/>	HOL02-DMZ-RT	<a href="#">rtb-0b8d284f56f330008</a>	-	-	No	<a href="#">vpc-0399cd2bbc16bd746   HOL...</a>	654654181466
<input type="checkbox"/>	HOL02-Production-RT	<a href="#">rtb-0ff8373ed5ca49679</a>	-	-	No	<a href="#">vpc-0399cd2bbc16bd746   HOL...</a>	654654181466
<input type="checkbox"/>	HOL02-Research-RT	<a href="#">rtb-0154ddb78120563d</a>	-	-	No	<a href="#">vpc-0399cd2bbc16bd746   HOL...</a>	654654181466

Figure 6

Afterward, we associate each route table with their subnet:

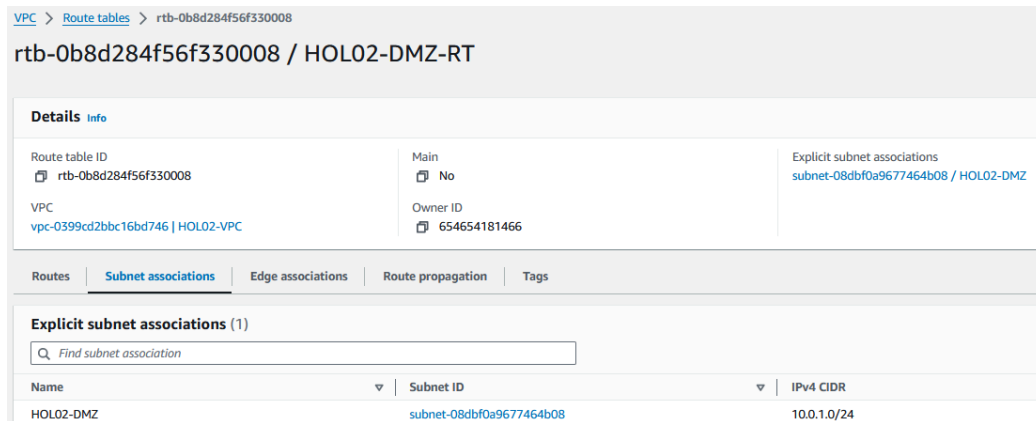


Figure 7

Moreover, the subnet that should be accessible from the internet is the one hosting the HOL02-VoilaServer, since this server is meant to share the team's findings with the public. Therefore, we must ensure that the subnet associated with the VoilaServer (the Production Subnet), has internet access.

To do that, first we create an internet gateway (*Figure 8*) and attach it to our VPC.




Figure 8- Internet gateway creation.

Then, we edit the route table associated with the Production Subnet (HOL02-Production) by adding a new route with 0.0.0.0/0 as the destination and the Internet Gateway as the target (*Figure 9*).

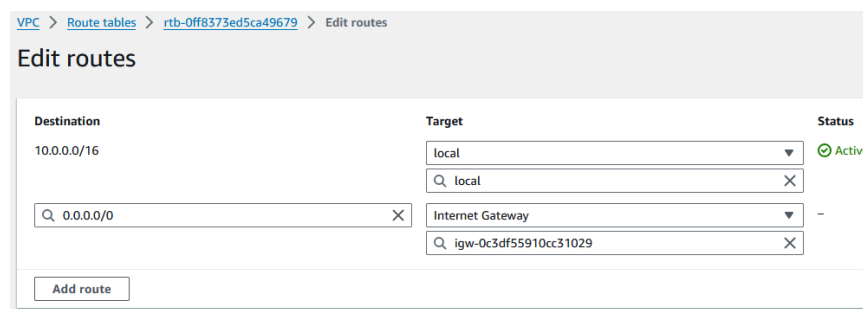


Figure 9

We follow the same steps for the route table associated with the HOL02-DMZ subnet. At the moment, the final network configuration is as follows:




Figure 10

## 2.3. Security groups

In AWS, security groups act as virtual firewalls for EC2 instances and other resources within the VPC. They basically control inbound and outbound traffic at the instance level, and this enables to specify which traffic is allowed to reach our instances.

- **Inbound Rules:** These rules define the traffic that is allowed to reach the instances.
- **Outbound Rules:** These rules define the outbound traffic that is allowed to leave the instances. For instance, we can specify that all outbound traffic is allowed, or restrict it to specific protocols and ports.

Still within the VPC Dashboard, we click on "Security Groups" in the sidebar. We create 3 different security groups (HOL02-DMZ-SG, HOL02-Production-SG and HOL02-Research-SG) with the specified ports from the assignment instructions (*table 1*). The security groups are created in the same way as *Figure 8*:



VPC > Security Groups > Create security group

## Create security group Info

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.

---

### Basic details

Security group name Info  
  
Name cannot be edited after creation.

Description Info

VPC Info

---

### Inbound rules Info

Type <small>Info</small>	Protocol <small>Info</small>	Port range <small>Info</small>	Source <small>Info</small>
SSH	TCP	22	Anywhere-IPv4 <small>0.0.0.0/0</small>
HTTPS	TCP	443	Anywhere-IPv4 <small>0.0.0.0/0</small>
Custom TCP	TCP	943	Anywhere-IPv4 <small>0.0.0.0/0</small>
Custom TCP	TCP	945	Anywhere-IPv4 <small>0.0.0.0/0</small>
Custom UDP	UDP	1194	Anywhere-IPv4 <small>0.0.0.0/0</small>

Add rule

Figure 11

Security group	Inbound rules	Outbund rules
HOL02-DMZ-SG	22(SSH) 443 (HTTPS) 943 (TCP) 945 (TCP) 1194 (UDP)	All traffic
HOL02-Production-SG	80 (HTTP) 443 (HTTPS) 22 (SSH) - only from HOL02-DMZ-SG	All traffic
HOL02-Research-SG	22 (SSH) - only from HOL02-DMZ-SG  8888 (TCP) - only from HOL02-DMZ-SG	All traffic

Table 1

For the Production security group, we will be allowing HTTP and HTTPS traffic from anywhere, and SSH traffic only from the DMZ security group.

## 2.4. EC2 instances

An EC2 instance is a virtual server in Amazon's EC2 for running applications on the AWS infrastructure. When creating an instance, we are essentially renting a virtual server from AWS to run our own applications.

For this assignment, we have to create three EC2 instances: HOL02-VPN, HOL02-VoilaServer and HOL02-Jupyter.

We go to the EC2 dashboard and launch an instance. For each instance, we select the characteristics following *Table 2*:

<b>EC2 instance name</b>	<b>Subnet associated</b>	<b>AMI</b>	<b>Instance type</b>
HOL02-VPN	HOL02-DMZ	Ubuntu 22.04 LTS	t2.micro
HOL02-VoilaServer	HOL02-Production	Amazon Linux 2	t2.micro
HOL02-Jupyter	HOL02-Research	Amazon Linux 2	t2.micro

*Table 2*

Here is an example of the EC2 instance configuration for HOL02-VPN:

**Name and tags** [Info](#)

Name

HOL02-VPN

Add additional tags

**▼ Application and OS Images (Amazon Machine Image)** [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Recents

Quick Start

Amazon Linux

aws

macOS

Mac

Ubuntu

ubuntu

Windows

Microsoft

Red Hat

Red Hat

SUSE Linux

SUSE

Browse more AMIs

Including AMIs from AWS, Marketplace and the Community

**Amazon Machine Image (AMI)**

Ubuntu Server 22.04 LTS (HVM), SSD Volume Type

Free tier eligible

ami-07d9b9ddc6cd8dd30 (64-bit (x86)) / ami-0568072f574d822a4 (64-bit (Arm))

Virtualization: hvm    ENA enabled: true    Root device type: ebs

**Description**

Canonical, Ubuntu, 22.04 LTS, amd64 jammy image build on 2024-02-07

Architecture

AMI ID

64-bit (x86)

ami-07d9b9ddc6cd8dd30

Verified provider

**▼ Key pair (login)** [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

HOL02

Create new key pair

**▼ Network settings** [Info](#)

VPC - required [Info](#)

vpc-0399cd2bbc16bd746 (HOL02-VPC)

10.0.0.0/16

Subnet [Info](#)

subnet-08dbf0a9677464b08

HOL02-DMZ

VPC: vpc-0399cd2bbc16bd746    Owner: 654654181466

Availability Zone: us-east-1d    IP addresses available: 251    CIDR: 10.0.1.0/24

Create new subnet

Auto-assign public IP [Info](#)

Disable

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

Common security groups [Info](#)

Select security groups

HOL02-DMZ-SG sg-03487e041144f48aa

Compare security group rules

Figure 12

9

HOL02-VPN requires an elastic IP, as specified in the assignment, so we have to go to the "Elastic IPs" section in the EC2 dashboard, create an Elastic IP and associate it with the EC2 instance HOL02-VPN:

EC2 > Elastic IP addresses > 35.169.222.223

35.169.222.223

Actions Associate Elastic IP address

**Summary**

Allocated IPv4 address 35.169.222.223	Type Public IP	Allocation ID eipalloc-01e03c24e821a9408	Reverse DNS record -
Association ID eipassoc-0207bbe3855763c54	Scope VPC	Associated instance ID i-0e3639d2321abdffa	Private IP address 10.0.1.201
Network interface ID eni-049dcf3edfda795b7	Network interface owner account ID 654654181466	Public DNS -	NAT Gateway ID -
Address pool Amazon	Network border group us-east-1		

Figure 13.-Elastic IP creation

EC2 > Elastic IP addresses > 35.169.222.223 > Associate Elastic IP address

**Associate Elastic IP address**

Choose the instance or network interface to associate to this Elastic IP address (35.169.222.223)

**Elastic IP address: 35.169.222.223**

**Resource type**  
Choose the type of resource with which to associate the Elastic IP address.

☒ Instance  
☐ Network interface

**Warning:** If you associate an Elastic IP address with an instance that already has an Elastic IP address associated, the previously associated Elastic IP address will be disassociated, but the address will still be allocated to your account. [Learn more](#)

If no private IP address is specified, the Elastic IP address will be associated with the primary private IP address.

**Instance**  
i-0e3639d2321abdffa

**Private IP address**  
The private IP address with which to associate the Elastic IP address.  
Choose a private IP address

**Reassociation**  
Specify whether the Elastic IP address can be reassociated with a different resource if it already associated with a resource.  
☐ Allow this Elastic IP address to be reassociated

Cancel Associate

Figure 14.- We associate the elastic IP to the instance.

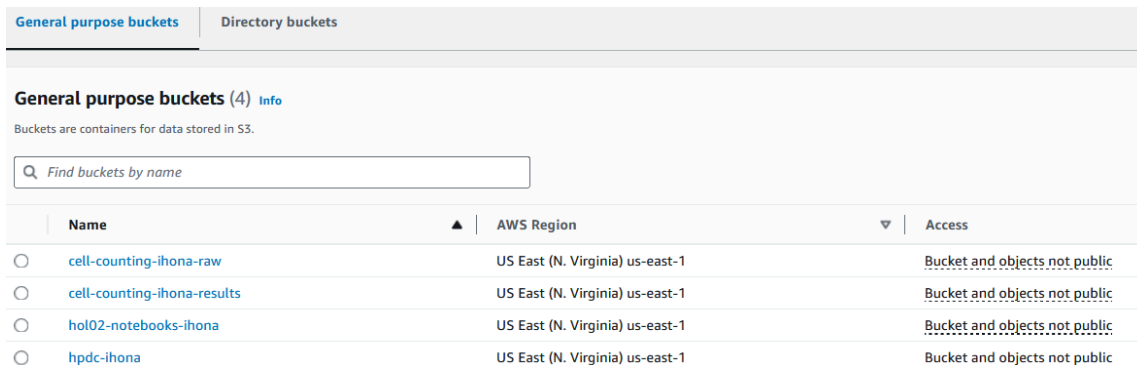
Here is an overview of the instances created:

Instances (4) Info											
Find instance by attribute or tag (case-sensitive)											
Any state											
Instance state: running Clear filters											
Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP	IPv6 IPs	Monitoring
Session02-Jup...	i-075555525b5237e2	Running	t2.micro	2/2 checks passed	View alarms	us-east-1b	ec2-54-210-243.com...	3.84.210.243	-	-	disabled
HOL02-VPN	i-0e3639d2321abdffa	Running	t2.micro	2/2 checks passed	View alarms	us-east-1d	-	35.169.222.223	35.169.222.223	-	disabled
HOL02-Jupyter	i-044f564095cfa0926	Running	t2.micro	2/2 checks passed	View alarms	us-east-1d	-	-	-	-	disabled
HOL02-VollaS...	i-0b17f644545b9263e	Running	t2.micro	2/2 checks passed	View alarms	us-east-1d	-	3.81.250.52	-	-	disabled
Security group name											
launch-wizard-1											
HOL02-DMZ-SG											
HOL02-Research-SG											
HOL02-Production-SG											

Figure 15

## 2.5. S3 buckets

Finally, we must create the S3 bucket that will include all the notebooks used by the research team. We navigate to the S3 dashboard, select S3 and click on “create bucket”. The bucket is named “hol02-notebooks-ihona”.



General purpose buckets		Directory buckets
<b>General purpose buckets (4)</b> <a href="#">Info</a>		
Buckets are containers for data stored in S3.		
<input type="text" value="Find buckets by name"/>		
Name	AWS Region	Access
<input type="radio"/> cell-counting-ihona-raw	US East (N. Virginia) us-east-1	<a href="#">Bucket and objects not public</a>
<input type="radio"/> cell-counting-ihona-results	US East (N. Virginia) us-east-1	<a href="#">Bucket and objects not public</a>
<input type="radio"/> hol02-notebooks-ihona	US East (N. Virginia) us-east-1	<a href="#">Bucket and objects not public</a>
<input type="radio"/> hpd-ihona	US East (N. Virginia) us-east-1	<a href="#">Bucket and objects not public</a>

Figure 16

## 3. OPEN VPN CONFIGURATION

OpenVPN is a robust and configurable VPN solution that allows to securely connect devices to a remote network over the internet. In this context, OpenVPN allows us to securely access the Jupyter server, which is not directly exposed to the public internet, by routing the connection through the VPN.

Firstly, we must connect to the EC2 instance named HOL02-VPN. To do it, we open a terminal in the computer and go to the directory where there is our private key. Then, we execute this command,

```
ssh -i aws01 ubuntu@35.169.222.223
```

where “aws01” is the name of the private key and 35.169.222.223 is the public IP of the instance. Once connected to the EC2 instance, we use the following commands to install OpenVPN on the EC2 instance:

```
ubuntu@ip-10-0-1-201:~$ sudo apt update -y
```

```

sudo apt update -y
sudo apt install ca-certificates gnupg wget net-tools -y
sudo wget https://as-repository.openvpn.net/as-repo-public.asc -qO
/etc/apt/trusted.gpg.d/as-repo-public.asc
sudo echo "deb [arch=amd64 signed-by=/etc/apt/trusted.gpg.d/as-repo-public.asc]
http://as-repository.openvpn.net/as/debian jammy main" | sudo tee
/etc/apt/sources.list.d/openvpn-as-repo.list
sudo apt update && sudo apt install openvpn-as -y

```

```

+++++
Access Server 2.13.1 has been successfully installed in /usr/local/openvpn_as
Configuration log file has been written to /usr/local/openvpn_as/init.log

Access Server Web UIs are available here:
Admin UI: https://10.0.1.201:943/admin
Client UI: https://10.0.1.201:943/
To login please use the "openvpn" account with [REDACTED] password.
(password can be changed on Admin UI)
+++++

```

The previous commands install OpenVPN Access Server, which is a package that simplifies the setup and management of OpenVPN connections.

The next step is to configure the OpenVPN. For that, we open a browser and go to this direction:

*<https://35.169.222.223:943/admin/>*

This redirects us to this page and we log in with the default username and password obtained from the terminal:




Figure 17

After the log in, we can access the OpenVPN web interface, so we navigate to Network Settings and change the hostname to the elastic IP as follows:

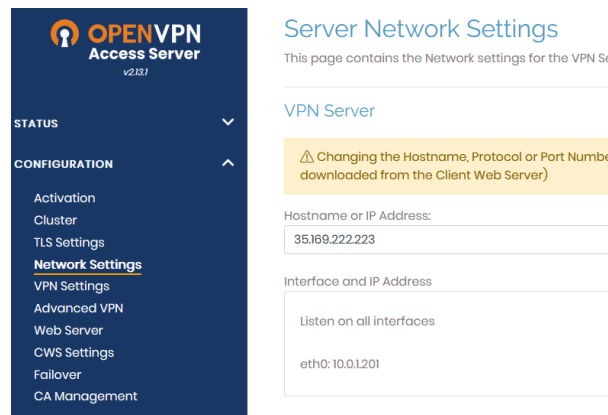


Figure 18

Then, we navigate to VPN Settings and add the subnets CIDR one per line:

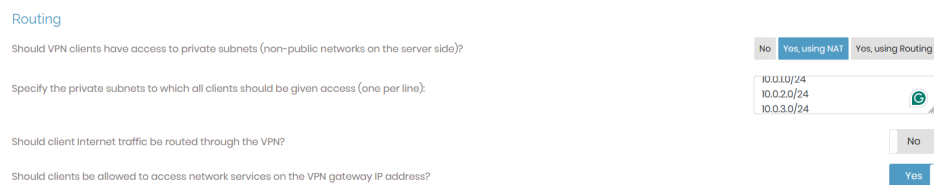


Figure 19

After configuring OpenVPN, we have to install the OpenVPN client. To achieve this, we open the browser and go to this direction:

<https://35.169.222.223:943>

We install the Yourself (user-locked) profile (Figure 16):




Figure 20

Now, if we provide the password obtained from the terminal, we are able to connect to the VPN (*Figure 17*). It will be necessary to be connected to the VPN in order to access the Jupyter server (port 8888 and port 22).

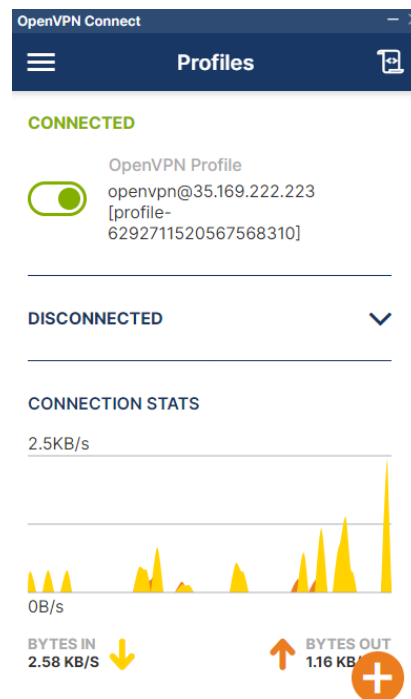


Figure 21- VPN connected

## 4. JUPYTER NOTEBOOK SETUP

First, we connect to the HOL02-Jupyter EC2 instance. Since the connection is over the VPN, we have to put the private IP address (*10.0.3.152*):

```
ssh -i aws01 ec2-user@10.0.3.152
```

Afterwards, we create a new python environment and we automatically activate it with this command:

```
echo "source jupyter-env/bin/activate" >> ~/.bashrc
```

We create a directory for the notebooks and run the command to initialize jupyter:

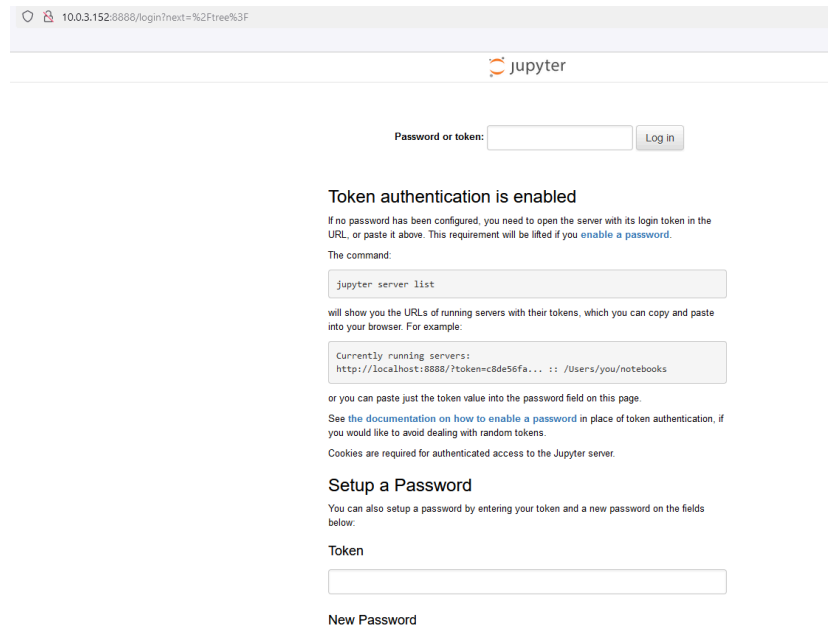
```
(jupyter-env) [ec2-user@ip-10-0-3-152 ~]$ mkdir notebooks
(jupyter-env) [ec2-user@ip-10-0-3-152 ~]$ cd notebooks
(jupyter-env) [ec2-user@ip-10-0-3-152 notebooks]$ jupyter notebook --ip=0.0.0.0
```

Finally, we go to the browser and paste this direction to be redirected to jupyter:



<http://10.0.3.152:8888/>

Once in the provided address, we must paste the token prompted in the terminal to access:



10.0.3.152:8888/login?next=%2Ftree%3F

jupyter

Password or token:  Log in

Token authentication is enabled

If no password has been configured, you need to open the server with its login token in the URL, or paste it above. This requirement will be lifted if you [enable a password](#).

The command:

```
jupyter server list
```

will show you the URLs of running servers with their tokens, which you can copy and paste into your browser. For example:

```
Currently running servers:
http://localhost:8888/?token=c8de5efa... :: /Users/you/notebooks
```

or you can paste just the token value into the password field on this page.

See [the documentation on how to enable a password](#) in place of token authentication, if you would like to avoid dealing with random tokens.

Cookies are required for authenticated access to the Jupyter server.

Setup a Password

You can also setup a password by entering your token and a new password on the fields below:

Token

New Password

Figure 22

For convenience, we can run this command to run Jupyter Notebook in the background and do not be asked for the token each time:

```
nohup jupyter notebook --ip=0.0.0.0 --notebook-dir=/home/ec2-user/notebooks
--no-browser > /home/ec2-user/notebooks/jupyter.log 2>&1 &
```

We create a sample Jupyter Notebook named “hol02-app.ipynb”.

It has been decided to use the file data.csv, which contains genomic data. This dataset is in an s3 bucket named hpdc-ihona. We download the dataset from the bucket by doing the following commands:

## 2. Import necessary libraries

```
[12]: import boto3
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from botocore import UNSIGNED
from botocore.client import Config
```

## 3. Import data from s3 bucket

```
[13]: session = boto3.Session()
s3 = boto3.client('s3',
region_name='us-east-1')
s3.download_file('hpdg-ihona', 'data.csv', 'data.csv')
df = pd.read_csv('data.csv')
```

## 4. Basic statistical analysis

```
[4]: df.head()
```

	ID	Chromosome	Position	Reference_Allele	Alternate_Allele	Genotype	Quality	Depth
0	1	1	100001	A	T	AA	30	50
1	2	1	150002	C	G	CC	28	45
2	3	2	500003	G	A	GG	35	60
3	4	2	750004	T	C	TT	32	55
4	5	3	300005	A	G	AG	25	40

```
[7]: df.shape
```

```
[7]: (10, 8)
```

Figure 23

After doing the analysis of the data (which contains some basic statistical analysis and visualizations), we can synchronize the jupyter notebook we just created with the s3 bucket created in the 2.5 section. We can do this either manually or with a tool named cron.

a) To do it manually, we have to run this command:

```
aws s3 sync /home/ec2-user/notebooks s3://hol02-notebooks-ihona --exclude "*.ipynb_checkpoints"
```

Where /home/ec2-user/notebooks/, is the path where the notebook is stored, and hol02-notebooks-ihona the bucket we want to save the jupyter notebook in:

```
(jupyter-env) [ec2-user@ip-10-0-3-152 notebooks]$ aws s3 sync /home/ec2-user/notebooks s3://hol02-notebooks-ihona --exclude "*.ipynb_checkpoints"
upload: ./data.csv to s3://hol02-notebooks-ihona/data.csv
upload: ./hol02-app.ipynb to s3://hol02-notebooks-ihona/hol02-app.ipynb
(jupyter-env) [ec2-user@ip-10-0-3-152 notebooks]$
```

We can check the bucket to ensure that everything has worked out fine:

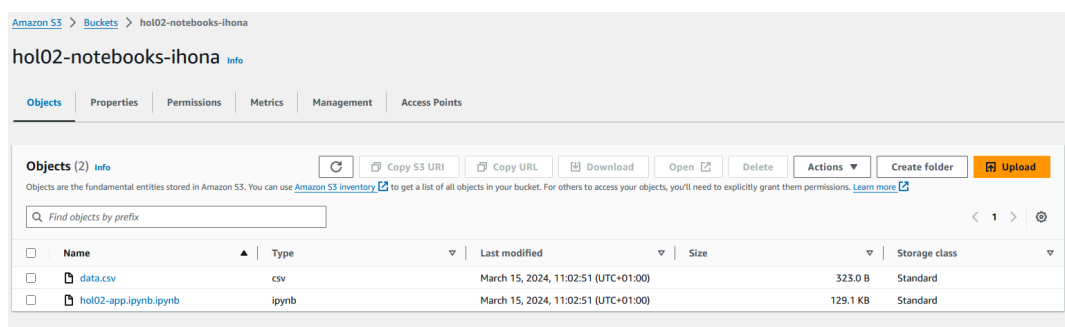


Figure 24

- b) Cron jobs are a powerful tool for scheduling tasks automatically at specific times. In order to do the synchronization with cron, first we have to install it using this command:

```
sudo yum install cronie -y
```

Then we start the cron service:

```
sudo systemctl start crond
```

And enable Cron Service to Start on Boot:

```
sudo systemctl enable crond
```

Then we execute this command:

```
echo -e "0 0 * * 0 sudo yum update -y\n0 3 * * * /usr/bin/aws s3 sync  
/home/ec2-user/notebooks s3://hol02-notebooks-ihona --exclude  
\"*.ipynb_checkpoints*\" | crontab -
```

The first cron job `0 0 * * 0 sudo yum update -y` is scheduled to run at 00:00 (midnight) on Sunday (day 0 of the week). This job will update the system's packages automatically using yum.

The second cron job `0 3 * * * /usr/bin/aws s3 sync /home/ec2-user/notebooks s3://hol02-notebooks-ihona --exclude "*.ipynb_checkpoints*"` is scheduled to run at 03:00 every day. It synchronizes the Jupyter notebooks from the specified directory (`/home/ec2-user/notebooks`) to the S3 bucket (`hol02-notebooks-ihona`), excluding any Jupyter checkpoint files.

We can verify the job scheduling with the command `crontab -l`

```
(jupyter-env) [ec2-user@ip-10-0-3-152 ~]$ echo -e "0 0 * * 0 sudo yum update -y\n0 3 * * * /usr/bin/aws s3 sync /home  
2-user/notebooks s3://hol02-notebooks-ihona --exclude \"*.ipynb_checkpoints*\" | crontab -  
(jupyter-env) [ec2-user@ip-10-0-3-152 ~]$ crontab -l  
0 0 * * 0 sudo yum update -y  
0 3 * * * /usr/bin/aws s3 sync /home/ec2-user/notebooks s3://hol02-notebooks-ihona --exclude \"*.ipynb_checkpoints*\"  
(jupyter-env) [ec2-user@ip-10-0-3-152 ~]$
```

This command lists all cron jobs for the current user. The output confirmed that both jobs are correctly scheduled.

## 5. VOILA SERVER CONFIGURATION

Voila turns Jupyter notebooks into standalone web applications, where the content of the Notebook's is displayed, with the option to hide the code. This is very useful for sharing data visualizations, reports, or interactive dashboards derived from Jupyter notebooks.

Nginx, on the other hand, is a high-performance web server software. In a real case, it could be used as the web server that hosts the Voila web application, making it accessible from the internet.

However, in order to simplify the activity, we will assume that the default page of nginx is the Voila server.

First of all, to configure the server, we connect through SSH to the EC2 instance named HOL02-VoilaServer, also accessible via its private IP.

Then, we use the following commands to install Nginx on the EC2 instance:

```
sudo dnf install nginx
sudo systemctl start nginx
sudo systemctl enable nginx
```

```
Installed:
generic-logos-httpd-18.0.0-12.amzn2023.0.3.noarch      gperftools-libs-2.9.1-1.amzn2023.0.3.x86_64
libunwind-1.4.0-5.amzn2023.0.2.x86_64                 nginx-1:1.24.0-1.amzn2023.0.2.x86_64
nginx-core-1:1.24.0-1.amzn2023.0.2.x86_64             nginxfilesystem-1:1.24.0-1.amzn2023.0.2.noarch
nginx-mimetypes-2.1.49-3.amzn2023.0.3.noarch

Complete!
[ec2-user@ip-10-0-2-96 ~]$ sudo systemctl start nginx
[ec2-user@ip-10-0-2-96 ~]$ sudo systemctl enable nginx
Created symlink /etc/systemd/system/multi-user.target.wants/nginx.service → /usr/lib/systemd/system/nginx.service.
[ec2-user@ip-10-0-2-96 ~]$
```

We can check that we can access the Voila Server from the internet by going to the browser and using the public IP address of the HOL02-VoilaServer EC2 instance:

*http://3.81.250.52*

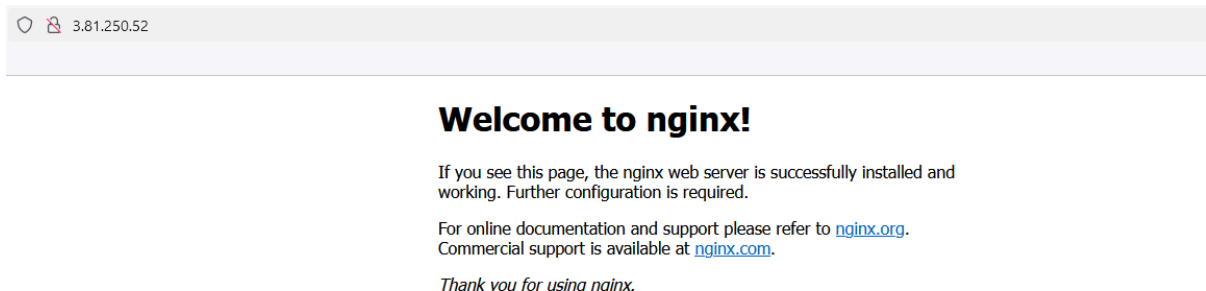


Figure 25

## 6. CHALLENGES AND TROUBLESHOOTING

During this assignment, we had to face several challenges and perform some troubleshooting to overcome them.

The first issue was regarding the connection to the HOL02-VPN EC2 instance. When trying to connect, we would get a time out message:

```
PS C:\Users\Ihona\AWS_Master> ssh -i aws01 ubuntu@35.169.222.223
ssh: connect to host 35.169.222.223 port 22: Connection timed out
```

To troubleshoot the problem, we had a look at the HOL02-DMZ subnet configuration and the correspondent route tables. We realized that the EC2 instance had an elastic IP attached but its route table was not linking to an internet gateway. If the subnet where the instance resides is not public or is not associated with a route table that has a route to the internet, the instance is not able to communicate to the internet or allow SSH access.

Therefore, we solved the issue by adding a new route with 0.0.0.0/0 as the destination and the Internet Gateway as the target.

Also, at the beginning, we were not able to connect to the VPN via OpenVPN client. The reason was that the connection was not taking the elastic IP so there was a connection time out. This issue was solved by downloading the user-locked profile.

Moreover, after connecting to HOL02-Jupyter EC2 instance, and creating and activating a virtual environment, we could not create a jupyter notebook. There was an error because we had not installed jupyter:

```
(jupyter-env) [ec2-user@ip-10-0-3-152 ~]$ jupyter notebook --ip=0.0.0.0
-bash: jupyter: command not found
```

However, we were not able to install jupyter because the HOL02-Research subnet, where the Jupyter instance is located, was effectively acting as a private subnet and did not have a direct route to the Internet Gateway. Therefore, the instance could not reach the internet directly to install packages. After doing some research, we found that to enable instances in private subnets access the internet (for software installation, updates, etc.) without being directly exposed to the internet, a **NAT Gateway** was used. This gateway is placed in a subnet that has access to the internet and allows outbound internet access for instances in private subnets.

We decided to create a NAT Gateway in the HOL02-Production subnet (*Figure 22*). Then, the route table for HOL02-Research subnet was updated to route outbound internet traffic (0.0.0.0/0) through this NAT Gateway (*Figure 23*).

**Create NAT gateway** [Info](#)

A highly available, managed Network Address Translation (NAT) service that instances in private subnets can use to connect to services in other VPCs, on-premises networks, or the internet.

**NAT gateway settings**

**Name - optional**  
Create a tag with a key of 'Name' and a value that you specify.  
HOL02-NAT  
The name can be up to 256 characters long.

**Subnet**  
Select a subnet in which to create the NAT gateway.  
subnet-0e0f5e298977dbc58 (HOL02-Production)

**Connectivity type**  
Select a connectivity type for the NAT gateway.  
☒ Public  
☐ Private

**Elastic IP allocation ID** [Info](#)  
Assign an Elastic IP address to the NAT gateway.  
eipalloc-02ffc99b6b7d6078a [Allocate Elastic IP](#)

[▶ Additional settings](#) [Info](#)

Figure 26

**Edit routes**

Destination	Target	Status	Propagated
10.0.0.0/16	local	Active	No
0.0.0.0/0	NAT Gateway	-	No

[Add route](#) [Cancel](#) [Preview](#) [Save changes](#)

Figure 27

**rtb-0154ddb78120563d / HOL02-Research-RT**

**Details** [Info](#)

Route table ID: rtb-0154ddb78120563d  
VPC: vpc-0399cd2bbc16bd746 | HOL02-VPC  
Main: No  
Owner ID: 654654181466  
Explicit subnet associations: subnet-05430e9d59aeecca / HOL02-Research

**Routes** [Subnet associations](#) [Edge associations](#) [Route propagation](#) [Tags](#)

**Routes (2)**

Destination	Target	Status
0.0.0.0/0	nat-0c3d2f09a48cf5f	Active
10.0.0.0/16	local	Active

Figure 28

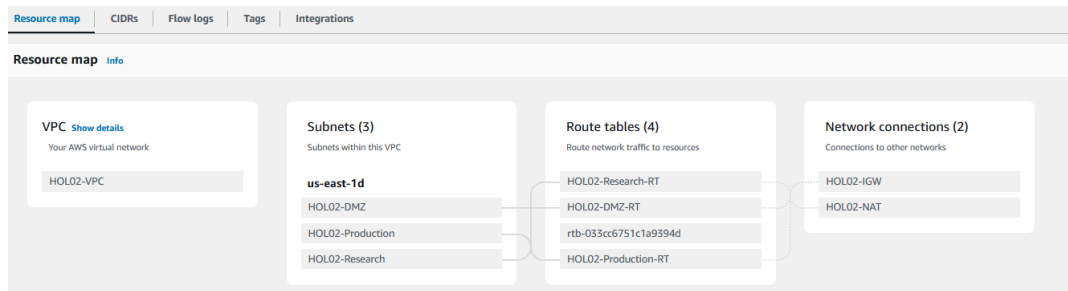


Figure 29.- Network configuration.

After doing these network configurations, we were able to successfully install jupyter in the research instance while maintaining a secure environment.

Lastly, there was an issue regarding the jupyter notebook used for the assignment. At first, it was thought that the notebook would get the data.csv file from the bucket hpdc-ihona. However, this made that for every time we wanted to run the notebook, we had to grant access to the virtual machine to access the bucket through the commands `aws configure` and `aws_session_token`. This was not very effective so it was decided to put the data directly in the bucket (hol02-notebooks-ihona) where the notebook is saved. To read the file, we only need to execute: `df = pd.read_csv('data.csv')`.

## 7. CONCLUSIONS

During this laboratory session, the requested infrastructure (VPC, subnets, route tables, security groups, EC2 instances and s3 buckets) was successfully set up. The OpenVPN was also successfully configured and was necessary in order to connect to the VoilaServer instance via SSH and the Jupyter instance. Moreover, a Jupyter notebook was set up in a safe environment only accessible through the VPN and also synchronized with the s3 bucket, both manually and through cron jobs.

On the other hand, the Voila Server was properly configured and the default Nginx page is accessible from the internet via the public IP, enabling the “production team” to share results with the public in a safe way.

Throughout the report, screenshots providing evidence, tests and troubleshooting have been provided.

In conclusion, AWS services were used to provide a solution and built an hybrid infrastructure for researchers in the cloud. This practical session has been very useful to learn new concepts regarding cloud computing with AWS and has undoubtedly improved my solvem-problem abilities.